

e-Tasmi AI Learning Loop

Tools, Services & Evidence Documentation

This document records the principal tools, services, APIs, development platforms, and supporting resources used to build and validate the e-Tasmi AI Learning Loop during the challenge.

Track 03 — Interactive Experiences & Learning Journeys for Islam

Individual entry

Documentation scope and evidence method

This booklet lets a reviewer see which external services support e-Tasmi, what each one is used for, and what the supplied screenshots actually show. It is an evidence record for the challenge submission, not a claim about commercial terms that were never on screen.

Categories covered

AI models and speech-to-text; Qur'an reference and library sources; the design, writing, and local development toolchain; and hosting, media, and session platforms.

How the evidence was collected

Eleven screenshots were supplied in the project folder `Evidence e-tasmi`. Each file was inspected before it was copied into `docs/evidence/screenshots/`. The original folder was not renamed or edited. Where a screen was too tall or mostly empty, the PDF shows the panel that carries the evidence, and the full image remains in the evidence folder.

Implementation details were checked against the project source and the challenge setup notes. Those checks are labelled as source or project records. They are not treated as something a screenshot proved by itself.

Evidence standard. Evidence statements are limited to what can be supported by the supplied screenshots and project records.

A screenshot proves the account or tool state that is visible on it. Source code and project configuration prove how the application is written. A screenshot does not establish a price, a plan name, or a licence unless that information is visible or is stated on the official page shown in the screenshot.

Four kinds of statement appear in this booklet:

- **Screenshot.** What is on the supplied image.
- **Project record.** What the challenge notes or example configuration document.
- **Source.** What the application code does.
- **Inference.** Not used to fill gaps. If a price is not shown, this booklet says so.

No credentials, API keys, passwords, or sensitive beneficiary data are included in this evidence package. The Cloudinary account email was covered on the documentation copy and marked as redacted. The original screenshot was left unchanged.

AI models and intelligent services

The learning loop separates three jobs. Quranpedia supplies the trusted Qur'an text. A Java comparison computes the word-level findings. The model explains those findings. The instructor accepts, edits, or rejects each finding before an ordinary result is published. The Qur'an text itself is not generated by the model.

OpenAI API

Purpose. Explanation of findings that have already been computed, and the existing Qur'an assistant on a separate model.

Implementation. Source: the explainer defaults to model id `gpt-6.1-so1` with reasoning effort high. That id can be overridden. The Qur'an assistant defaults to `gpt-4o-mini`. A legacy path on `gpt-4o` is not the learning-loop path. EV-01 does not show these request settings.

Evidence. The Home screen introduces GPT-6.1 Sol and records 123,378 tokens, 84 responses and chat completions, a 22.1% prompt-cache hit rate, and 84 requests. October spend is labeled Personal and shows \$0.68 of \$100.00, with a credit balance of \$2.83. No plan name or unit price is shown, and the totals do not prove that every request was the explainer.




Figure EV-01 — OpenAI API Home, with GPT-6.1 Sol announced and project usage recorded, including an October spend card and a credit balance. Evidence purpose: supports documentation of the OpenAI API environment used during challenge development. It does not show the application request body.

ElevenLabs

Purpose. Speech-to-text for a recitation recording, before the deterministic comparison.

Implementation. Source: when the speech provider is ElevenLabs, the client uses model id `scribe_v2`. The challenge setup notes document that choice. If the provider is unset, the selector uses OpenAI transcription instead. The live environment file is not reproduced here.

Evidence. EV-02 is a signed-in Speech to Text page with a card for Scribe Realtime v2. The transcript titles are not identified as e-Tasmi recordings. Pricing status not established by the supplied evidence.




Figure EV-02 — ElevenLabs Speech to Text, including a Scribe Realtime v2 card. Evidence purpose: supports documentation of ElevenLabs access during challenge development. The model id used by the application is established from source, not from this screen. No price is visible.

Qur'an and Islamic content sources

Two sources have different jobs. Quranpedia is the trusted text for the learning-loop comparison. Quran Foundation is the Qur'an library used for content such as verses, translations, tafsir, and audio. Neither source is model output. The comparison runs in Java against the trusted reference, and the explainer is instructed that the word-level comparison is already done.

Quranpedia API

Role. Source: the reference client defaults to `https://api.quranpedia.net/v1` and does not send credentials. Publication of a normal analysis requires that reference source. The comparison label in the service is the Quranpedia Hafs text.

Evidence. EV-04 is the public documentation page. The page states free, read-only access, no authentication, and the base URL above. The same page states attribution rules for republished dumps and rate limits of 120 requests per minute and 10,000 per day. The figure shows the upper portion; the full page is in the evidence folder.

Pricing. The supplied documentation page states that the API is free and that no authentication is required. That statement is from the page, not from an invoice.




Figure EV-04 — Quranpedia documentation: free read-only access, no authentication, and the documented base URL. Evidence purpose: supports the reference source and the terms printed on that page. Rate limits appear further down the full screenshot.

Quran Foundation

Role. Source: the Content API client is configured from server environment variables and uses OAuth2 client credentials. It is the library source. It is not the text the learning loop compares against.

Evidence. EV-03 is the official documentation portal, showing Content APIs, a JavaScript SDK, and manual OAuth2. It does not show an e-Tasmi billing account or an approved production credential.

Pricing. Pricing status not established by the supplied screenshot.




Figure EV-03 — Quran Foundation documentation portal. Evidence purpose: identifies the official Content API and OAuth2 documentation used as the library reference. It does not show account pricing.

Development and design toolchain

These tools supported design, writing, and local development. They are not the runtime services in the learning loop. ChatGPT, in particular, is a separate product from the OpenAI API billed on EV-01.

LaTeX and TeXstudio

Role. Authoring the challenge presentation. EV-06 shows TeXstudio with the deck source, XeLaTeX selected, and a preview of pages 1 to 3 of 28.

Licensing. The screenshot does not show a price. LaTeX is distributed by the LaTeX Project as free software. TeXstudio's own licence is not shown on this screen.




Figure EV-06 — TeXstudio and XeLaTeX open on the e-Tasmi challenge presentation. Evidence purpose: shows the documentation toolchain used for the submission deck.

ChatGPT

Role. Research and ideation. EV-05 is a project named AI Challenge. The visible thread discusses attempt progression using the existing attempt fields. It is not an API usage dashboard and it is not evidence of API billing.

Pricing. Pricing status not established by the supplied evidence.




Figure EV-05 — ChatGPT project used while shaping the learning journey. Evidence purpose: distinguishes the chat product, used for development discussion, from the OpenAI API used at runtime.

Figma and Git

EV-07 shows a Figma team workspace and starter templates. The file visible on that screen was edited two years earlier. It proves account access. It does not prove a challenge-period design file for this interface, and it does not show a price. The application is maintained in Git. No Git hosting screenshot was supplied, so no Git plan is stated. EV-07 is retained in [screenshots/figma/](#).

Railway

The diagram illustrates the data flow between two components. At the bottom, a box for 'e-tasmi AI Learning Loop' (e-tasmi.com) is shown with a green 'Online' status. An upward-pointing arrow connects it to a box for 'MySQL' database, which also shows a green 'Online' status. The MySQL box includes a sub-label 'mysql-volume'.

Docker

docker desktop

PERSONAL

Drive

🔊

🔒

🔧

⚙️

⋮

Sign in

🏠

Gordon

📦

Containers

🖼️

Images

💾

Volumes

🔑

Kubernetes

🏗️

Builds

📁

Models

🔧

MCP Toolkit

🌐

Docker Hub

📄

Logs

🔍

Extensions

Containers

View feedback

🔍 Search

Container CPU usage

12.20% / 1600% (16 CPUs available)

Container memory usage

3.95GB / 7.39GB

Show charts

🔍 Search

🔍 Only show running containers

☐	Name	Container ID	Image	Port(s)	CPU (%)	Memory usage	Memory (%)	Disk res	Actions
☒	docker	-	-	-	0.73%	410.9MB / 15.1%	5.3%	124.3M	
☒	etasm	-	-	-	1.05%	665.22MB / 30.2%	8.58%	118.59B	
☒	etasm-chal	-	-	-	0.86%	783.35MB / 30.2%	10.1%	118.09B	
☒	etasm-chalic	9255e85014fc	etasm-chal: 8084.9080		0.16%	255.98MB / 7.57%	3.3%	7.09MB	
☐	etasm-chalic	b3ac3dc3c3eb	caddy:2-chal: 8444.443		0%	35.03MB / 7.57%	0.45%	42.9MB	
☐	etasm-chalic	b1867233051c	mysql:8.4: 3317233051		0.7%	458.1MB / 7.57%	6.04%	43.4MB	
☐	etasm-chalic	afc5528a373d	phpmyadmin: 8083.80		0%	24.32MB / 7.57%	0.31%	24.5MB	

Zoom and Cloudinary

A screenshot of the Zoom Open API Call Logs page. The top navigation bar includes the Zoom logo and "Marketplace". On the left sidebar, under the "Developer" section, there are links for "Create apps", "API call logs" (which is highlighted), "Billing", "Community", and "Documentation". The main content area has a header with "API Call Logs" and "Webhook Logs - Legacy Webhook Logs". Below this, a message states: "Below are the Zoom Open API call logs generated by the app you have authorized. Only the calls generated within the last 30 days are listed below." There is a table with two columns: "Name" and "Type". A single row is visible with a blue icon containing a plus sign and the word "Special" next to it in the "Name" column, and "Private_Media" in the "Type" column.

athensv1

[Users](#)
[Groups](#)

Billing

Plan Details

Billing Information

Invoices

Product environment settings

API Keys

Upload

1 users (limit 3)

[Last login](#)
[Status](#)

Name	Status	Product Environments	Groups	Last Login	Two-Factor
Athenas Abdothaman Edit Reset	active	all	0 groups	Oct 05, 2026	on

[Add](#)
[Invite](#)

Page 1 of 1

[Previous](#)
[Next](#)

e-Tasmi AI Learning Loop

Licensing, pricing, and evidence register

A blank price is recorded as not established. The only free-of-charge statement below is the one printed on the Quranpedia documentation page in EV-04. Dollar amounts for OpenAI are the figures visible on EV-01, not a reconstructed invoice.

Service	Purpose and evidence	Pricing / access note
OpenAI API	Explainer for compared findings. EV-01: GPT-6.1 Sol banner and project usage. Source: default model <code>gpt-6.1-sol</code> , reasoning effort high. Assistant default <code>gpt-4o-mini</code> .	October spend card labeled Personal, \$0.68 of \$100.00. Credit balance \$2.83. Plan name and unit price not shown.
ElevenLabs	Speech-to-text when selected. EV-02: Speech to Text page and Scribe Realtime v2 card. Source model id <code>scribe_v2</code> .	Pricing status not established by the supplied evidence.
Quranpedia	Trusted comparison text. EV-04 and source base URL https://api.quranpedia.net/v1 .	The documentation page states that the API is free and needs no authentication. Rate limits are stated on that page.
Quran Foundation	Qur'an library, not the comparison text. EV-03: official docs portal. Source uses OAuth2 client credentials.	Pricing status not established by the supplied screenshot.
Zoom	Tasmi session meetings. EV-11: app <code>e-tasmi</code> , Private_OAuth.	Pricing status not established by the supplied evidence.
Cloudinary	Media handling in source. EV-10: active account, displayed user limit 3, Upgrade control. Email redacted.	Pricing status not established by the supplied evidence.
Railway	Application hosting. EV-09: app and MySQL online at <code>e-tasmi.com</code> .	Pricing status not established by the supplied evidence.
Docker	Local challenge environment. EV-08: Desktop header PERSONAL; MySQL 8.4 and the app container running.	PERSONAL is visible. No price is shown.
Figma	Design account. EV-07: team and templates. Visible file edited two years earlier.	Pricing status not established by the supplied evidence.
LaTeX / TeXstudio	Presentation authoring. EV-06: TeXstudio and XeLaTeX on the challenge deck.	No price on the screenshot. LaTeX Project distributes LaTeX as free software.
ChatGPT	Development discussion. EV-05: project AI Challenge. Not the runtime API.	Pricing status not established by the supplied evidence. Not an API invoice.

References and evidence index

Official destinations checked for this booklet. Quran Foundation is cited at its current documentation portal, which is the portal shown in EV-03 and the address already used in the project source.

OpenAI API	https://platform.openai.com/
ElevenLabs	https://elevenlabs.io/
Quran Foundation	https://api-docs.quran.foundation/
Quranpedia	https://quranpedia.net/api-docs
Zoom Developer	https://developers.zoom.us/
Cloudinary	https://cloudinary.com/
Railway	https://railway.com/
Figma	https://www.figma.com/
Docker	https://www.docker.com/
LaTeX Project	https://www.latex-project.org/

Screenshot index

EV-01	OpenAI API usage and GPT-6.1 Sol	Page 3
EV-02	ElevenLabs Speech to Text	Page 3
EV-03	Quran Foundation documentation	Page 4
EV-04	Quranpedia API documentation	Page 4
EV-05	ChatGPT, AI Challenge project	Page 5
EV-06	TeXstudio / XeLaTeX presentation	Page 5
EV-07	Figma team (file retained; not a primary figure)	Page 5
EV-08	Docker Desktop, challenge containers	Page 6
EV-09	Railway application and MySQL	Page 6
EV-10	Cloudinary account, email redacted	Page 6
EV-11	Zoom developer app e-tasmi	Page 6

Implementation notes checked in source

The explainer defaults to `gpt-6.1-sol` with reasoning effort high, and it explains findings already produced by `RecitationComparisonEngine`. ElevenLabs `scribe_v2` is the speech-to-text model id when that provider is selected. Quranpedia is the trusted reference required for publication of a normal analysis. Pending findings block that publication until the instructor accepts, edits, or rejects them. The stack is Java servlets and JSP, JDBC, and MySQL.

No credentials, API keys, passwords, or sensitive beneficiary data are included in this evidence package. Originals remain in Evidence `e-tasmi`. Organized copies, including the redacted Cloudinary image, are in `docs/evidence/screenshots/`. The full register is `docs/evidence/evidence-index.md`.